

```
In [308]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import warnings
import seaborn as sns
warnings.filterwarnings('ignore')
plt.style.use('fivethirtyeight')
%matplotlib inline
pd.set_option('display.max_columns', 26)
```

```
In [309]: df = pd.read_csv('kidney_disease.csv')
df.head()
```

Out[309]:

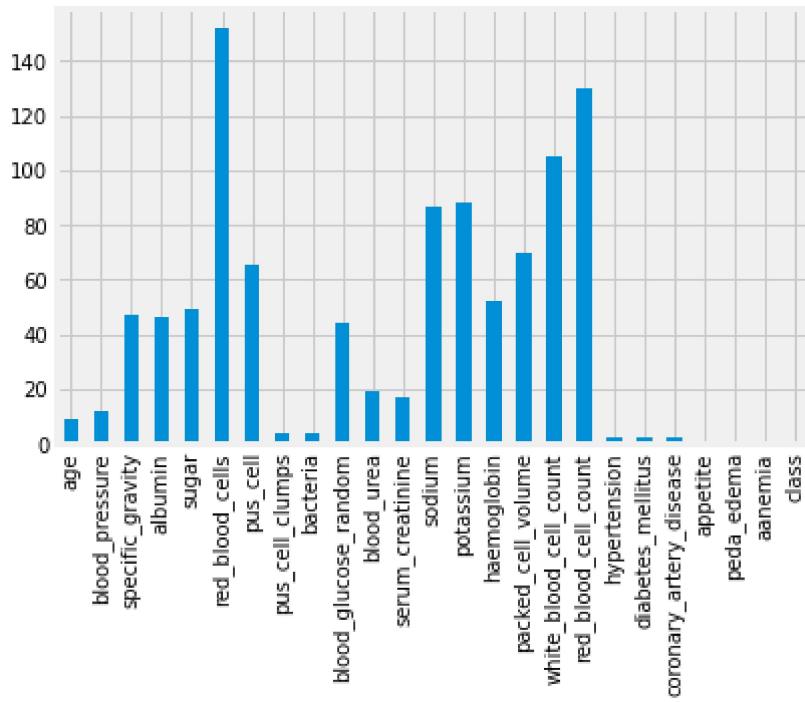
| | id | age | bp | sg | al | su | rbc | pc | pcc | ba | bgr | bu | sc | soc |
|---|----|------|------|-------|-----|-----|--------|----------|------------|------------|-------|------|-----|-------|
| 0 | 0 | 48.0 | 80.0 | 1.020 | 1.0 | 0.0 | NaN | normal | notpresent | notpresent | 121.0 | 36.0 | 1.2 | NaN |
| 1 | 1 | 7.0 | 50.0 | 1.020 | 4.0 | 0.0 | NaN | normal | notpresent | notpresent | NaN | 18.0 | 0.8 | NaN |
| 2 | 2 | 62.0 | 80.0 | 1.010 | 2.0 | 3.0 | normal | normal | notpresent | notpresent | 423.0 | 53.0 | 1.8 | NaN |
| 3 | 3 | 48.0 | 70.0 | 1.005 | 4.0 | 0.0 | normal | abnormal | present | notpresent | 117.0 | 56.0 | 3.8 | 111.0 |
| 4 | 4 | 51.0 | 80.0 | 1.010 | 2.0 | 0.0 | normal | normal | notpresent | notpresent | 106.0 | 26.0 | 1.4 | NaN |



```
In [310]: # dropping id column
df.drop('id', axis = 1, inplace = True)
```

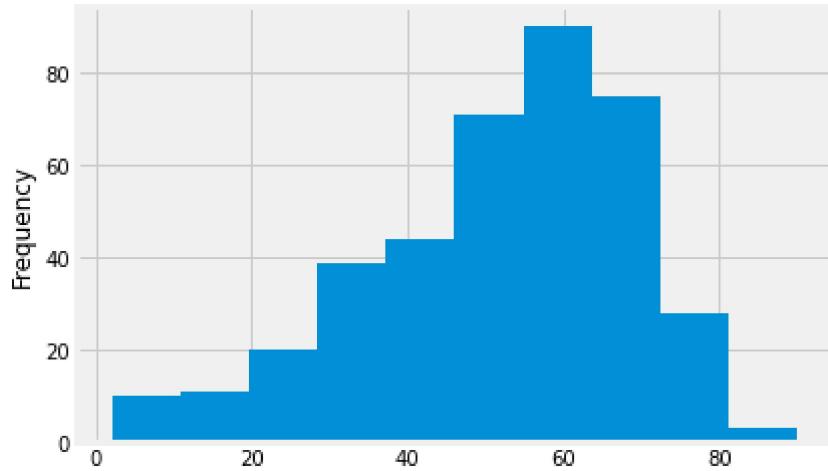
```
In [311]: #rename
df.columns = ['age', 'blood_pressure', 'specific_gravity', 'albumin', 'sugar',
              'pus_cell_clumps', 'bacteria', 'blood_glucose_random', 'blood_ur',
              'potassium', 'haemoglobin', 'packed_cell_volume', 'white_blood_c',
              'hypertension', 'diabetes_mellitus', 'coronary_artery_disease',
              'aanemia', 'class']
```

```
In [312]: # Here we are plotting the graph to see the null values in the dataset.  
p = df.isnull().sum().plot.bar()
```



```
In [313]: #checking the distribution of the age column  
df['age'].plot(kind='hist')
```

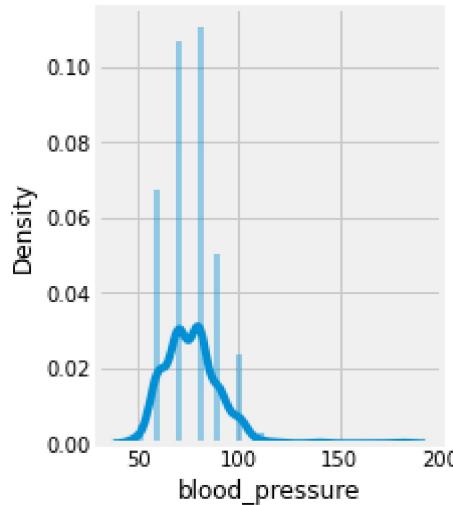
Out[313]: <AxesSubplot:ylabel='Frequency'>



In [314]: # distribution of bp column

```
plt.subplot(121), sns.distplot(df['blood_pressure'])
plt.show
```

Out[314]: <function matplotlib.pyplot.show(close=None, block=None)>



In [315]: df.head()

Out[315]:

| | age | blood_pressure | specific_gravity | albumin | sugar | red_blood_cells | pus_cell | pus_cell_cl |
|---|------|----------------|------------------|---------|-------|-----------------|----------|-------------|
| 0 | 48.0 | 80.0 | 1.020 | 1.0 | 0.0 | NaN | normal | notp |
| 1 | 7.0 | 50.0 | 1.020 | 4.0 | 0.0 | NaN | normal | notp |
| 2 | 62.0 | 80.0 | 1.010 | 2.0 | 3.0 | normal | normal | notp |
| 3 | 48.0 | 70.0 | 1.005 | 4.0 | 0.0 | normal | abnormal | p |
| 4 | 51.0 | 80.0 | 1.010 | 2.0 | 0.0 | normal | normal | notp |

In [316]: df.describe()

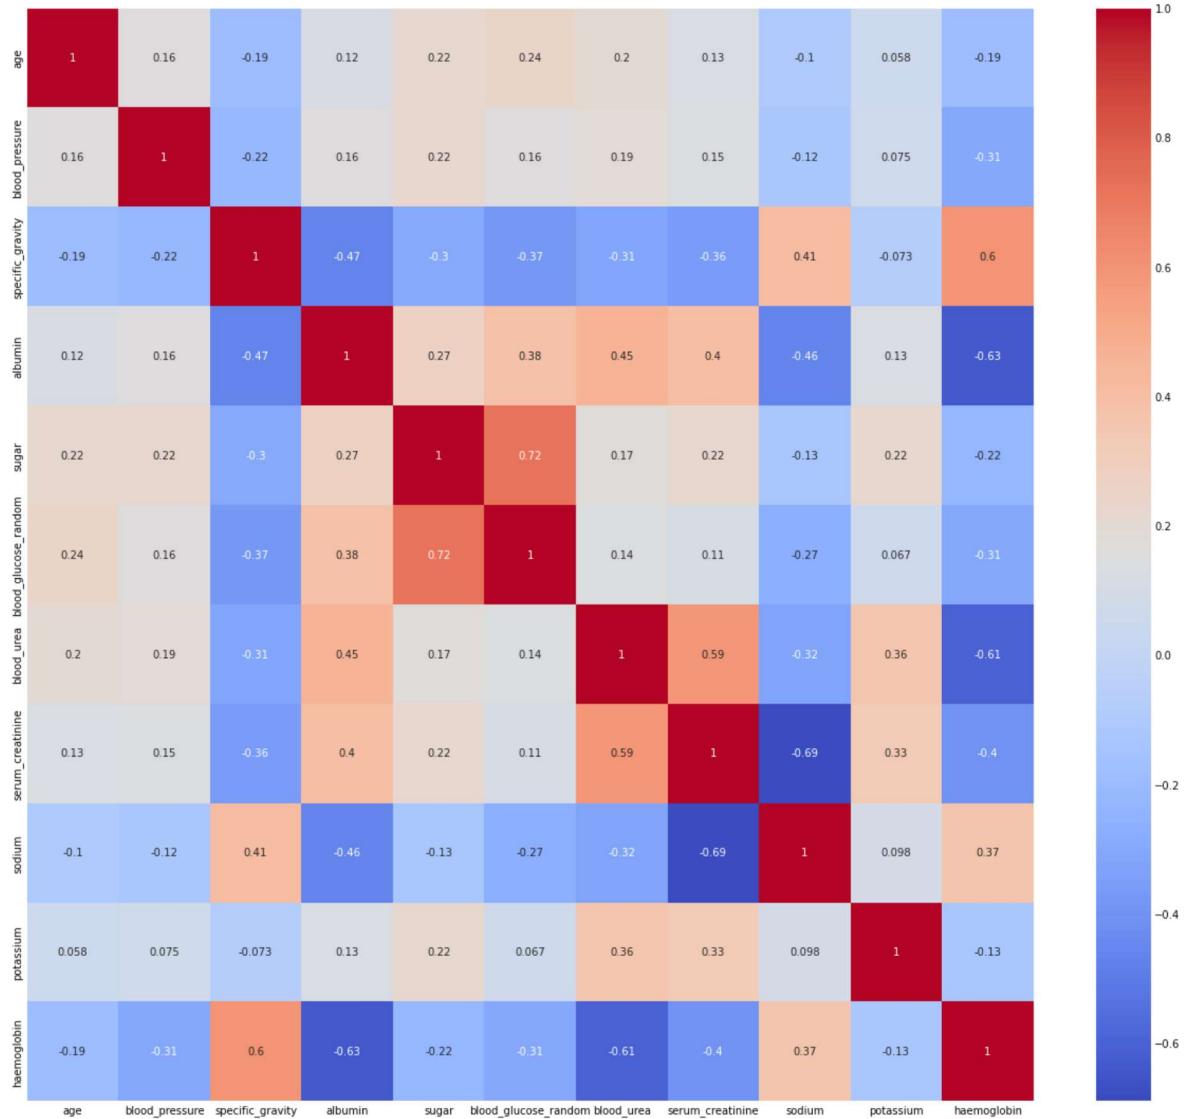
Out[316]:

| | age | blood_pressure | specific_gravity | albumin | sugar | blood_glucose_rand |
|-------|------------|----------------|------------------|------------|------------|--------------------|
| count | 391.000000 | 388.000000 | 353.000000 | 354.000000 | 351.000000 | 356.000 |
| mean | 51.483376 | 76.469072 | 1.017408 | 1.016949 | 0.450142 | 148.036 |
| std | 17.169714 | 13.683637 | 0.005717 | 1.352679 | 1.099191 | 79.281 |
| min | 2.000000 | 50.000000 | 1.005000 | 0.000000 | 0.000000 | 22.000 |
| 25% | 42.000000 | 70.000000 | 1.010000 | 0.000000 | 0.000000 | 99.000 |
| 50% | 55.000000 | 80.000000 | 1.020000 | 0.000000 | 0.000000 | 121.000 |
| 75% | 64.500000 | 80.000000 | 1.020000 | 2.000000 | 0.000000 | 163.000 |
| max | 90.000000 | 180.000000 | 1.025000 | 5.000000 | 5.000000 | 490.000 |

In [317]: #Finding the Correlation between the plots

```
plt.figure(figsize = (19,19))
sns.heatmap(df.corr(), annot = True, cmap = 'coolwarm') # Looking for strong correlations
```

Out[317]: <AxesSubplot:>



In [318]: df['packed_cell_volume'] = pd.to_numeric(df['packed_cell_volume'], errors='coerce')
df['white_blood_cell_count'] = pd.to_numeric(df['white_blood_cell_count'], errors='coerce')
df['red_blood_cell_count'] = pd.to_numeric(df['red_blood_cell_count'], errors='coerce')

In [319]: # Extracting categorical and numerical columns

```
cat_cols = [col for col in df.columns if df[col].dtype == 'object']
num_cols = [col for col in df.columns if df[col].dtype != 'object']
```

```
In [320]: # Looking at unique values in categorical columns
for col in cat_cols:
    print(f"{col} has {df[col].unique()} values\n")
```

```
red_blood_cells has [nan 'normal' 'abnormal'] values
pus_cell has ['normal' 'abnormal' nan] values
pus_cell_clumps has ['notpresent' 'present' nan] values
bacteria has ['notpresent' 'present' nan] values
hypertension has ['yes' 'no' nan] values
diabetes_mellitus has ['yes' 'no' ' yes' '\tno' '\tyes' nan] values
coronary_artery_disease has ['no' 'yes' '\tno' nan] values
appetite has ['good' 'poor' nan] values
peda_edema has ['no' 'yes' nan] values
aanemia has ['no' 'yes' nan] values
class has ['ckd' 'ckd\t' 'notckd'] values
```

```
In [321]: df['diabetes_mellitus'].replace(to_replace = {'\tno':'no', '\tyes':'yes', ' yes': 'yes'}, inplace=True)
df['coronary_artery_disease'] = df['coronary_artery_disease'].replace(to_replace = {'\tno': 'no', '\tno': 'no', '\tno': 'no'}, inplace=True)
df['class'] = df['class'].replace(to_replace = {'ckd\t': 'ckd', 'notckd': 'notckd'}, inplace=True)
```

```
In [322]: df['class'] = df['class'].map({'ckd': 0, 'not ckd': 1})
df['class'] = pd.to_numeric(df['class'], errors='coerce')
```

```
In [323]: cols = ['diabetes_mellitus', 'coronary_artery_disease', 'class']
for col in cols:
    print(f"{col} has {df[col].unique()} values\n")
```

```
diabetes_mellitus has ['yes' 'no' nan] values
coronary_artery_disease has ['no' 'yes' nan] values
class has [0 1] values
```

```
In [324]: # checking for null values  
df.isna().sum().sort_values(ascending = False)
```

```
Out[324]: red_blood_cells      152  
red_blood_cell_count        131  
white_blood_cell_count     106  
potassium                  88  
sodium                     87  
packed_cell_volume          71  
pus_cell                   65  
haemoglobin                52  
sugar                      49  
specific_gravity            47  
albumin                    46  
blood_glucose_random       44  
blood_urea                  19  
serum_creatinine            17  
blood_pressure               12  
age                         9  
bacteria                   4  
pus_cell_clumps             4  
hypertension                 2  
diabetes_mellitus           2  
coronary_artery_disease     2  
appetite                    1  
peda_edema                 1  
aanemia                     1  
class                       0  
dtype: int64
```

```
In [325]: df[num_cols].isnull().sum()
```

```
Out[325]: age                  9  
blood_pressure                12  
specific_gravity              47  
albumin                     46  
sugar                        49  
blood_glucose_random          44  
blood_urea                    19  
serum_creatinine              17  
sodium                       87  
potassium                    88  
haemoglobin                  52  
packed_cell_volume             71  
white_blood_cell_count        106  
red_blood_cell_count          131  
dtype: int64
```

```
In [326]: df[cat_cols].isnull().sum()
```

```
Out[326]: red_blood_cells      152
          pus_cell             65
          pus_cell_clumps       4
          bacteria              4
          hypertension           2
          diabetes_mellitus      2
          coronary_artery_disease 2
          appetite               1
          peda_edema             1
          aanemia                1
          class                  0
          dtype: int64
```

```
In [327]: # filling null values, we will use two methods, random sampling for higher null
# mean/mode sampling for lower null values
```

```
def random_value_imputation(feature):
    random_sample = df[feature].dropna().sample(df[feature].isna().sum())
    random_sample.index = df[df[feature].isnull()].index
    df.loc[df[feature].isnull(), feature] = random_sample

def impute_mode(feature):
    mode = df[feature].mode()[0]
    df[feature] = df[feature].fillna(mode)
```

```
In [328]: # filling num_cols null values using random sampling method
```

```
for col in num_cols:
    random_value_imputation(col)
```

```
In [329]: df[num_cols].isnull().sum()
```

```
Out[329]: age                  0
          blood_pressure        0
          specific_gravity       0
          albumin                0
          sugar                  0
          blood_glucose_random   0
          blood_urea              0
          serum_creatinine        0
          sodium                 0
          potassium              0
          haemoglobin            0
          packed_cell_volume      0
          white_blood_cell_count  0
          red_blood_cell_count    0
          dtype: int64
```

```
In [330]: # filling "red_blood_cells" and "pus_cell" using random sampling method and re  
random_value_imputation('red_blood_cells')  
random_value_imputation('pus_cell')  
  
for col in cat_cols:  
    impute_mode(col)
```

```
In [331]: df[cat_cols].isnull().sum()
```

```
Out[331]: red_blood_cells      0  
pus_cell          0  
pus_cell_clumps  0  
bacteria         0  
hypertension      0  
diabetes_mellitus 0  
coronary_artery_disease 0  
appetite          0  
peda_edema        0  
aanemia           0  
class             0  
dtype: int64
```

```
In [332]: for col in cat_cols:  
    print(f"{col} has {df[col].nunique()} categories\n")
```

```
red_blood_cells has 2 categories  
  
pus_cell has 2 categories  
  
pus_cell_clumps has 2 categories  
  
bacteria has 2 categories  
  
hypertension has 2 categories  
  
diabetes_mellitus has 2 categories  
  
coronary_artery_disease has 2 categories  
  
appetite has 2 categories  
  
peda_edema has 2 categories  
  
aanemia has 2 categories  
  
class has 2 categories
```

```
In [333]: from sklearn.preprocessing import LabelEncoder  
  
le = LabelEncoder()  
  
for col in cat_cols:  
    df[col] = le.fit_transform(df[col])
```

```
In [334]: df.head()
```

Out[334]:

| | age | blood_pressure | specific_gravity | albumin | sugar | red_blood_cells | pus_cell | pus_cell_cl |
|---|------|----------------|------------------|---------|-------|-----------------|----------|-------------|
| 0 | 48.0 | 80.0 | 1.020 | 1.0 | 0.0 | | 1 | 1 |
| 1 | 7.0 | 50.0 | 1.020 | 4.0 | 0.0 | | 1 | 1 |
| 2 | 62.0 | 80.0 | 1.010 | 2.0 | 3.0 | | 1 | 1 |
| 3 | 48.0 | 70.0 | 1.005 | 4.0 | 0.0 | | 1 | 0 |
| 4 | 51.0 | 80.0 | 1.010 | 2.0 | 0.0 | | 1 | 1 |



```
In [335]: ind_col = [col for col in df.columns if col != 'class']  
dep_col = 'class'  
X = df[ind_col]  
y = df[dep_col]
```

```
In [336]: # splitting data into training and test set
```

```
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30, ra
```

Logistic Regression

```
In [337]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
import pandas as pd
# Load data into a Pandas DataFrame
data = pd.read_csv('kidney_disease.csv')
# Separate the input (X) and output (y) variables
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Create a Logistic Regression object and fit it to the training data
Classifier = LogisticRegression()
Classifier.fit(X_train, y_train)
# Predict the output values for the testing set
y_pred = Classifier.predict(X_test)
# Calculate the accuracy score
accuracy = accuracy_score(y_test, y_pred)
# Print the accuracy score
print("Accuracy:", accuracy)
```

Accuracy: 0.925

K-NEAREST NEIGHBOUR

```
In [338]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

knn = KNeighborsClassifier()
knn.fit(X_train, y_train)

# accuracy score, confusion matrix and classification report of knn

knn_acc = accuracy_score(y_test, knn.predict(X_test))

print(f"Training Accuracy of KNN is {accuracy_score(y_train, knn.predict(X_train))}")
print(f"Test Accuracy of KNN is {knn_acc} \n")

print(f"Confusion Matrix :- \n{confusion_matrix(y_test, knn.predict(X_test))}")
print(f"Classification Report :- \n {classification_report(y_test, knn.predict(X_test))}")
```

Training Accuracy of KNN is 0.8

Test Accuracy of KNN is 0.75

Confusion Matrix :-

```
[[38 14]
 [ 6 22]]
```

Classification Report :-

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.86 | 0.73 | 0.79 | 52 |
| 1 | 0.61 | 0.79 | 0.69 | 28 |
| accuracy | | | 0.75 | 80 |
| macro avg | 0.74 | 0.76 | 0.74 | 80 |
| weighted avg | 0.78 | 0.75 | 0.76 | 80 |

In []:

DECISION TREE CLASSIFIER

```
In [339]: from sklearn.tree import DecisionTreeClassifier

dtc = DecisionTreeClassifier()
dtc.fit(X_train, y_train)

# accuracy score, confusion matrix and classification report of decision tree

dtc_acc = accuracy_score(y_test, dtc.predict(X_test))

print(f"Training Accuracy of Decision Tree Classifier is {accuracy_score(y_train, dtc.predict(X_train))}")
print(f"Test Accuracy of Decision Tree Classifier is {dtc_acc} \n")

print(f"Confusion Matrix :- \n{confusion_matrix(y_test, dtc.predict(X_test))}")
print(f"Classification Report :- \n {classification_report(y_test, dtc.predict(X_test))}
```

Training Accuracy of Decision Tree Classifier is 1.0
Test Accuracy of Decision Tree Classifier is 0.9625

Confusion Matrix :-

```
[[52  0]
 [ 3 25]]
```

Classification Report :-

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.95 | 1.00 | 0.97 | 52 |
| 1 | 1.00 | 0.89 | 0.94 | 28 |
| accuracy | | | 0.96 | 80 |
| macro avg | 0.97 | 0.95 | 0.96 | 80 |
| weighted avg | 0.96 | 0.96 | 0.96 | 80 |

RANDOM FOREST CLASSIFIER

```
In [340]: from sklearn.ensemble import RandomForestClassifier

rd_clf = RandomForestClassifier(criterion = 'entropy', max_depth = 11, max_features = 3)
rd_clf.fit(X_train, y_train)

# accuracy score, confusion matrix and classification report of random forest

rd_clf_acc = accuracy_score(y_test, rd_clf.predict(X_test))

print(f"Training Accuracy of Random Forest Classifier is {accuracy_score(y_train, rd_clf.predict(X_train))}")
print(f"Test Accuracy of Random Forest Classifier is {rd_clf_acc} \n")

print(f"Confusion Matrix :- \n{confusion_matrix(y_test, rd_clf.predict(X_test))}")
print(f"Classification Report :- \n {classification_report(y_test, rd_clf.predict(X_test))}")
```



Training Accuracy of Random Forest Classifier is 1.0
Test Accuracy of Random Forest Classifier is 0.9625

Confusion Matrix :-

```
[[52  0]
 [ 3 25]]
```

Classification Report :-

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.95 | 1.00 | 0.97 | 52 |
| 1 | 1.00 | 0.89 | 0.94 | 28 |
| accuracy | | | 0.96 | 80 |
| macro avg | 0.97 | 0.95 | 0.96 | 80 |
| weighted avg | 0.96 | 0.96 | 0.96 | 80 |

SVM

```
In [341]: from sklearn.svm import SVC
classifier =SVC(random_state = 0)
classifier.fit(X_train, y_train)
print(accuracy)
```

0.925

LINEAR REGRESSION

```
In [342]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split
import pandas as pd

# Separate the input (X) and output (y) variables
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create a Linear Regression object and fit it to the training data
regressor = LinearRegression()
regressor.fit(X_train, y_train)

# Predict the output values for the testing set
y_pred = regressor.predict(X_test)

# Calculate the R-squared score
accuracy = r2_score(y_test, y_pred)

# Print the accuracy score
print("Accuracy (R-squared):", accuracy)
```

Accuracy (R-squared): 0.6679328207916022

RANDOM FOREST REGRESSION

```
In [343]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
#Instantiate and fit the RandomForestClassifier
forest = RandomForestClassifier()
forest.fit(X_train, y_train)
# Make predictions for the test set
y_pred_test = forest.predict(X_test)
# View accuracy score
accuracy_score(y_test, y_pred_test)
# View confusion matrix for test data and predictions
confusion_matrix(y_test, y_pred_test)
# View the classification report for test data and predictions
print(classification_report(y_test, y_pred_test))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.95 | 1.00 | 0.97 | 52 |
| 1 | 1.00 | 0.89 | 0.94 | 28 |
| accuracy | | | 0.96 | 80 |
| macro avg | 0.97 | 0.95 | 0.96 | 80 |
| weighted avg | 0.96 | 0.96 | 0.96 | 80 |

POLYNOMIAL REGRESSION

```
In [344]: from sklearn.preprocessing import PolynomialFeatures
import numpy as np
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
# Transform the input data to include polynomial features up to degree 2
poly = PolynomialFeatures(degree=2)
x_poly = poly.fit_transform(X)
lin_reg_2 = LinearRegression()
lin_reg_2.fit(x_poly, y)
```

Out[344]: LinearRegression()

```
In [345]: lr = LinearRegression()
lr.fit(X_train, y_train)
y_pred = lr.predict(X_test)
print(r2_score(y_test, y_pred))
```

0.6679328207916022

Naive bayes

```
In [346]: X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size = 0.2,random_state=0)
# training the model on training set
from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb.fit(X_train, y_train)

# making predictions on the testing set
y_pred = gnb.predict(X_test)

# comparing actual response values (y_test) with predicted response values (y_pred)
from sklearn import metrics
nav_acc=metrics.accuracy_score(y_test, y_pred)*100
print("Gaussian Naive Bayes model accuracy(in %):", nav_acc)
```

Gaussian Naive Bayes model accuracy(in %): 95.0

kernal svm

```
In [347]: # Import the Libraries
import numpy as np
import matplotlib.pyplot as plt
from sklearn import svm, datasets

X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size = 0.2,random_state=0)
from sklearn.svm import SVC
svclassifier = SVC(kernel='rbf')
svclassifier.fit(X_train, y_train)
y_pred = svclassifier.predict(X_test)
from sklearn.metrics import classification_report, confusion_matrix
print("KERNEL SVM WITH GAUSSIAN KERNEL")
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
from sklearn.svm import SVC
svclassifier = SVC(kernel='sigmoid')
svclassifier.fit(X_train, y_train)
y_pred = svclassifier.predict(X_test)
from sklearn.metrics import classification_report, confusion_matrix
print("KERNEL SVM WITH SIGMOID KERNEL")
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

KERNEL SVM WITH GAUSSIAN KERNEL

```
[[52  0]
 [28  0]]
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.65 | 1.00 | 0.79 | 52 |
| 1 | 0.00 | 0.00 | 0.00 | 28 |
| accuracy | | | 0.65 | 80 |
| macro avg | 0.33 | 0.50 | 0.39 | 80 |
| weighted avg | 0.42 | 0.65 | 0.51 | 80 |

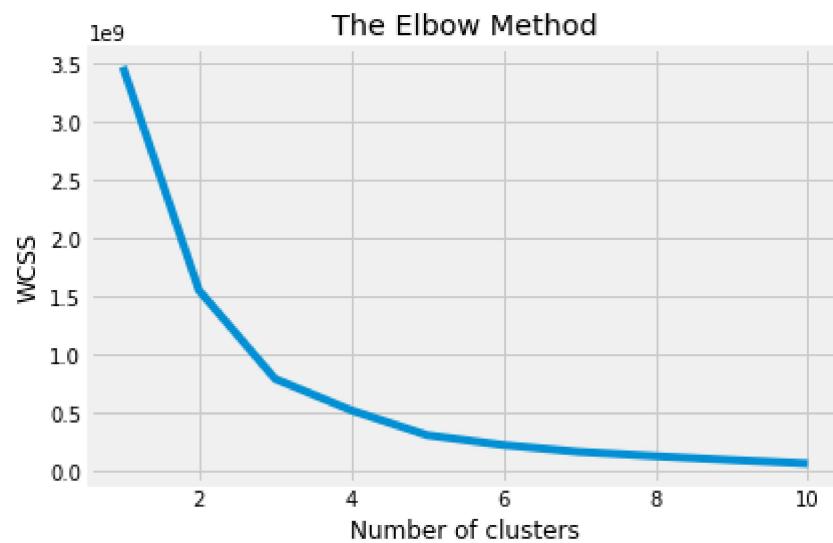
KERNEL SVM WITH SIGMOID KERNEL

```
[[31 21]
 [25  3]]
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.55 | 0.60 | 0.57 | 52 |
| 1 | 0.12 | 0.11 | 0.12 | 28 |
| accuracy | | | 0.42 | 80 |
| macro avg | 0.34 | 0.35 | 0.34 | 80 |
| weighted avg | 0.40 | 0.42 | 0.41 | 80 |

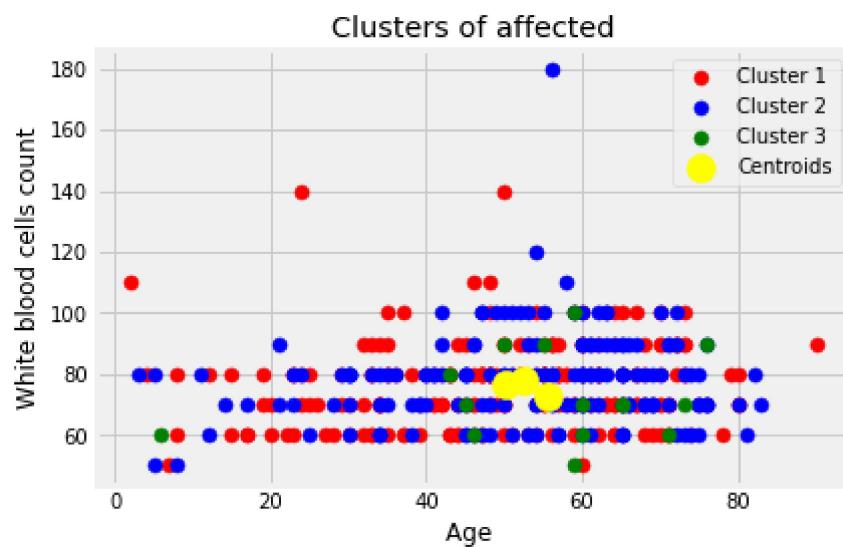
K mean clustering

```
In [348]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
#dataset
df = pd.read_csv('kidney_disease.csv')
x = df.iloc[:,[1,17]].values
#importing k means
from sklearn.cluster import KMeans
wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters = i, init = 'k-means++', random_state = 42)
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)
plt.plot(range(1, 11), wcss)
plt.title('The Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.show()
```



```
In [349]: kmeans = KMeans(n_clusters = 3, init = 'k-means++', random_state = 42)
y_kmeans = kmeans.fit_predict(X)

# Visualising the clusters
plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s = 50, c = 'red', label = 'Cluster 1')
plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s = 50, c = 'blue', label = 'Cluster 2')
plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], s = 50, c = 'green', label = 'Cluster 3')
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1], s = 100, c = 'yellow', label = 'Centroids')
plt.title('Clusters of affected')
plt.xlabel('Age ')
plt.ylabel('White blood cells count')
plt.legend()
plt.show()
```



HIERARCHICAL CLUSTERING

```
In [ ]: import scipy.cluster.hierarchy as sch
dendrogram = sch.dendrogram(sch.linkage(X, method = 'ward'))
plt.title('Dendrogram')
plt.xlabel('hemoglobin')
plt.ylabel('Euclidean distance')
plt.show()
```

```
In [ ]: from sklearn.cluster import AgglomerativeClustering
hc = AgglomerativeClustering(n_clusters = 3, affinity = 'euclidean', linkage = 'ward')
y_hc = hc.fit_predict(X)
```

```
In [ ]: plt.scatter(X[y_hc == 0, 0], X[y_hc == 0, 1], s = 50, c = 'green', label = 'C0')
plt.scatter(X[y_hc == 1, 0], X[y_hc == 1, 1], s = 50, c = 'magenta', label = 'C1')
plt.scatter(X[y_hc == 2, 0], X[y_hc == 2, 1], s = 50, c = 'orange', label = 'C2')

plt.title('Clusters of affected')
plt.xlabel('white blood cells')
plt.ylabel('Age')
plt.legend()
plt.show()
```

Decision tree Regression

```
In [ ]: from sklearn.metrics import mean_squared_error,r2_score
```

```
In [262]: from math import sqrt
```

```
In [263]: mse=mean_squared_error(y_test,y_pred)
```

```
In [264]: print(mse)
```

```
0.6125
```

```
In [265]: rmse=sqrt(mean_squared_error(y_test,y_pred))
```

```
In [266]: print(rmse)
```

```
0.7826237921249264
```

SVR

```
In [267]: from sklearn.svm import SVR
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size = 0.2,random_state=42)

# Fit the SVR model to the training data
clf = SVR(kernel='rbf', C=1e3, gamma=0.1)
clf.fit(X_train, y_train)

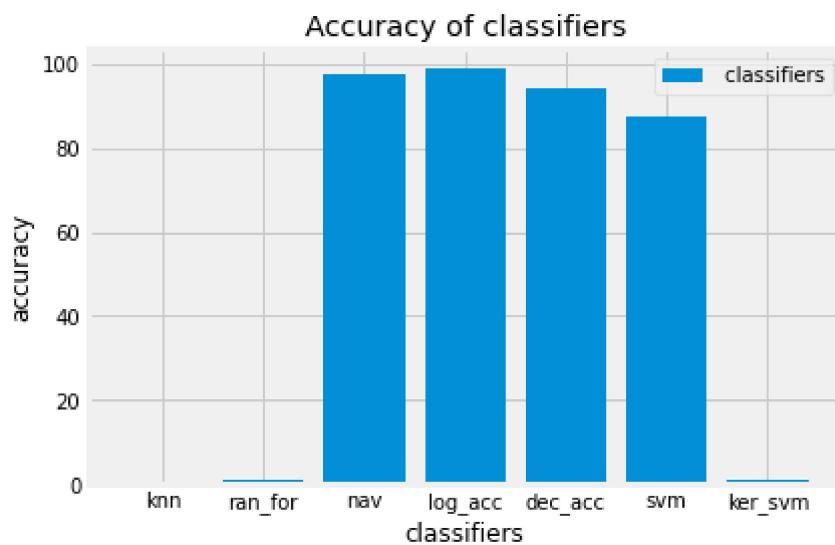
# Predict the target values for the test data using the trained model
y_pred = clf.predict(X_test)

# Evaluate the performance of the model on the test data using mean squared error
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error:", mse)
```

Mean Squared Error: 0.23052499998807502

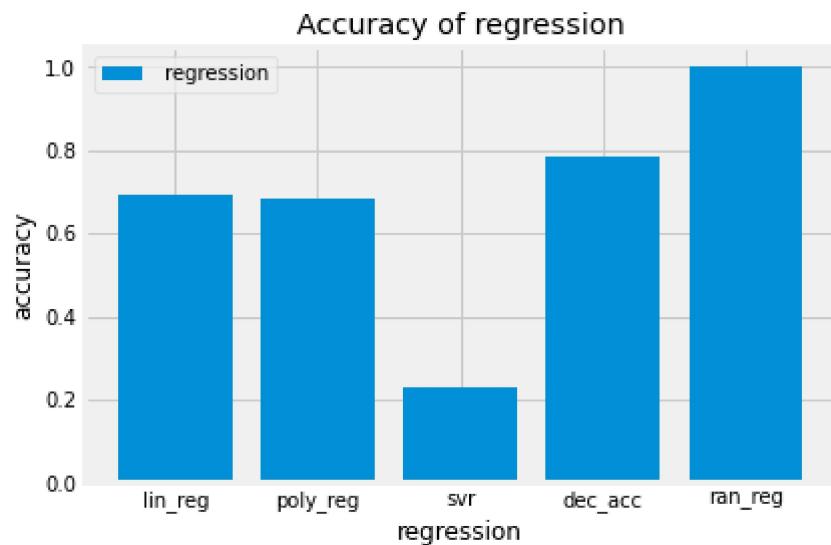
```
In [268]: x=['knn','ran_for','nav','log_acc','dec_acc','svm','ker_svm']
y=knn_acc,rd_clf_acc,nav_acc,98.7,94.1,87.5,1

X_axis = np.arange(len(x))
plt.bar(X_axis , y, label = ' classifiers')
plt.xticks(X_axis, x)
plt.xlabel("classifiers")
plt.ylabel("accuracy")
plt.title("Accuracy of classifiers")
plt.legend()
plt.show()
```



```
In [269]: x=['lin_reg','poly_reg','svr','dec_acc','ran_reg']
y=accuracy,0.682,mse,rmse,1

X_axis = np.arange(len(x))
plt.bar(X_axis , y, label = ' regression')
plt.xticks(X_axis, x)
plt.xlabel(" regression")
plt.ylabel("accuracy")
plt.title("Accuracy of regression")
plt.legend()
plt.show()
```



```
In [ ]:
```