

```
In [149]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import warnings
import seaborn as sns
warnings.filterwarnings('ignore')
plt.style.use('fivethirtyeight')
%matplotlib inline
pd.set_option('display.max_columns', 26)
```

```
In [150]: df = pd.read_csv('kidney_disease.csv')
df.head()
```

Out[150]:

	id	age	bp	sg	al	su	rbc	pc	pcc	ba	bgr	bu	sc	sod	pot	hei
0	0	48.0	80.0	1.020	1.0	0.0	NaN	normal	notpresent	notpresent	121.0	36.0	1.2	NaN	NaN	1
1	1	7.0	50.0	1.020	4.0	0.0	NaN	normal	notpresent	notpresent	NaN	18.0	0.8	NaN	NaN	1
2	2	62.0	80.0	1.010	2.0	3.0	normal	normal	notpresent	notpresent	423.0	53.0	1.8	NaN	NaN	
3	3	48.0	70.0	1.005	4.0	0.0	normal	abnormal	present	notpresent	117.0	56.0	3.8	111.0	2.5	1
4	4	51.0	80.0	1.010	2.0	0.0	normal	normal	notpresent	notpresent	106.0	26.0	1.4	NaN	NaN	1

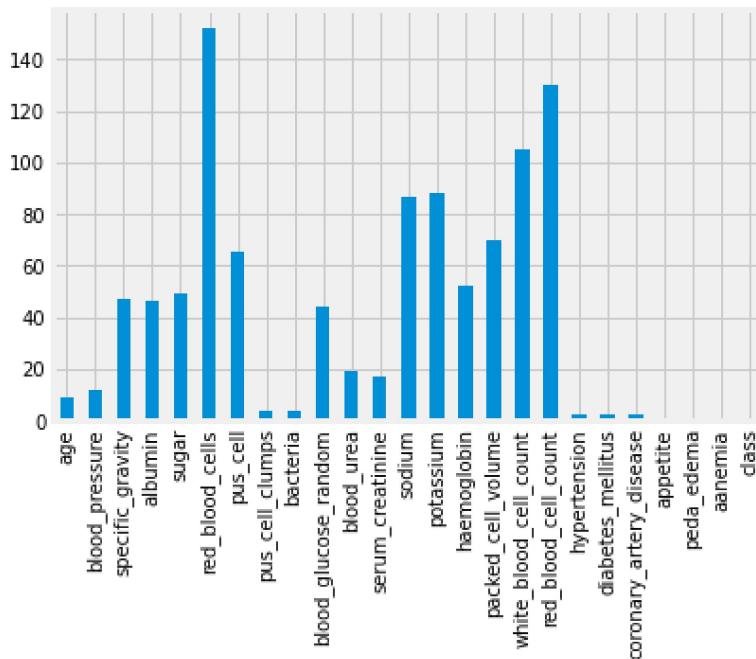


```
In [151]: # dropping id column
df.drop('id', axis = 1, inplace = True)
```

In [152]: #rename

```
df.columns = ['age', 'blood_pressure', 'specific_gravity', 'albumin', 'sugar', 'red_blood_cell', 'pus_cell_clumps', 'bacteria', 'blood_glucose_random', 'blood_urea', 'semen', 'potassium', 'haemoglobin', 'packed_cell_volume', 'white_blood_cell_count', 'hypertension', 'diabetes_mellitus', 'coronary_artery_disease', 'appetite', 'anaemia', 'class']
```

In [153]: # Here we are plotting the graph to see the null values in the dataset.
`p = df.isnull().sum().plot.bar()`



In [154]: `df.head()`

Out[154]:

	age	blood_pressure	specific_gravity	albumin	sugar	red_blood_cells	pus_cell	pus_cell_clumps	b
0	48.0	80.0	1.020	1.0	0.0	NaN	normal	notpresent	notp
1	7.0	50.0	1.020	4.0	0.0	NaN	normal	notpresent	notp
2	62.0	80.0	1.010	2.0	3.0	normal	normal	notpresent	notp
3	48.0	70.0	1.005	4.0	0.0	normal	abnormal	present	notp
4	51.0	80.0	1.010	2.0	0.0	normal	normal	notpresent	notp

In [155]: df.describe()

Out[155]:

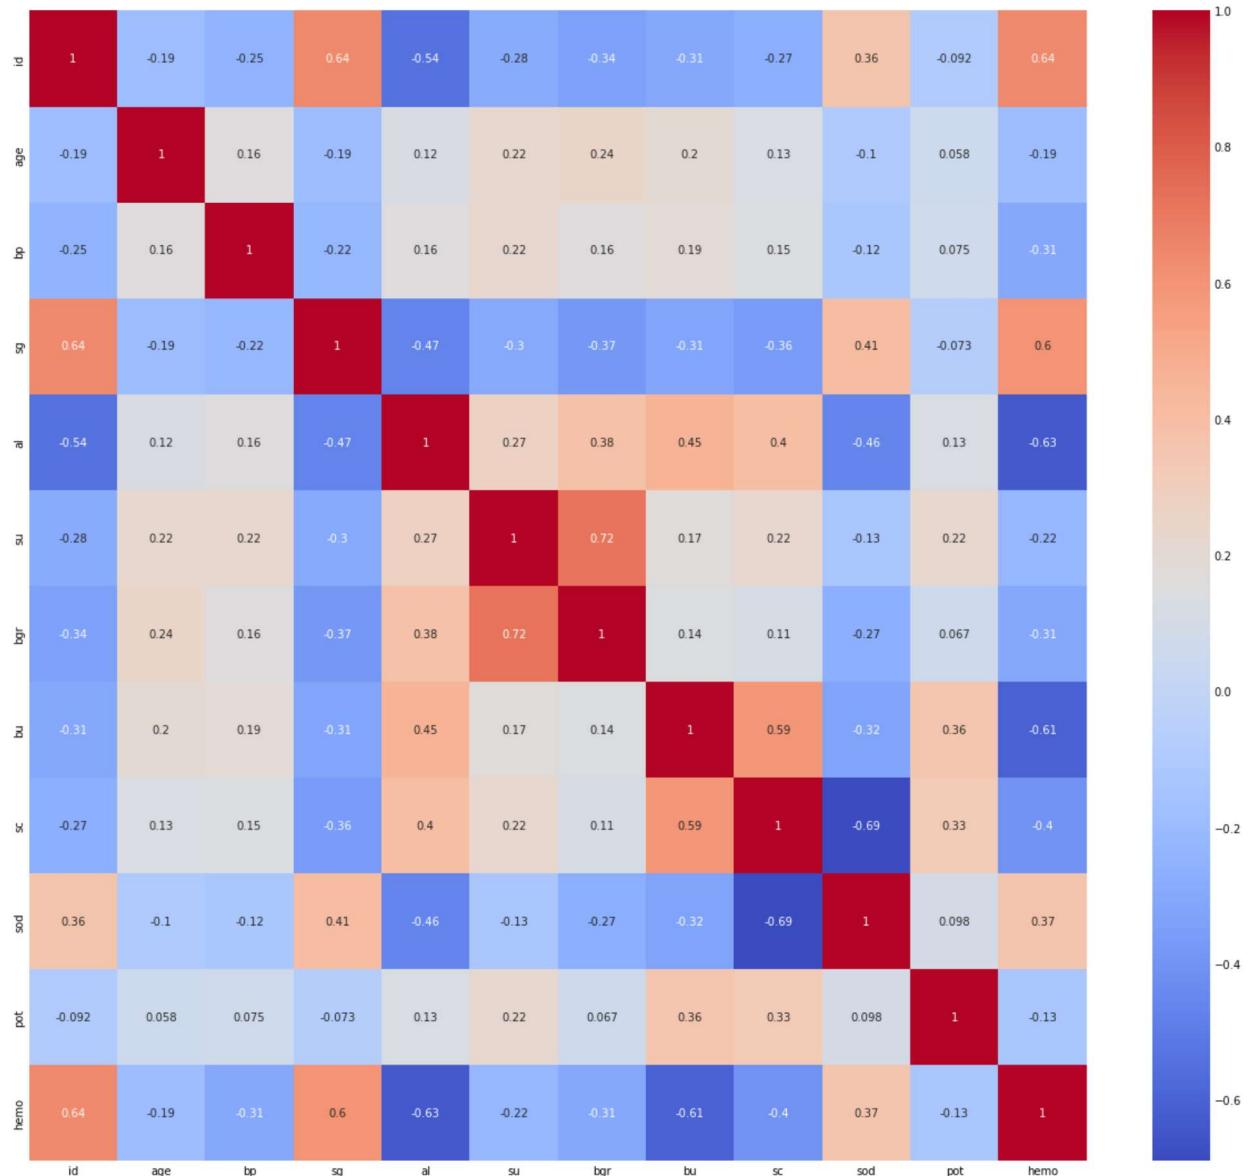
	age	blood_pressure	specific_gravity	albumin	sugar	blood_glucose_random	blood_
count	391.000000	388.000000	353.000000	354.000000	351.000000	356.000000	381.000000
mean	51.483376	76.469072	1.017408	1.016949	0.450142	148.036517	57.42
std	17.169714	13.683637	0.005717	1.352679	1.099191	79.281714	50.50
min	2.000000	50.000000	1.005000	0.000000	0.000000	22.000000	1.50
25%	42.000000	70.000000	1.010000	0.000000	0.000000	99.000000	27.00
50%	55.000000	80.000000	1.020000	0.000000	0.000000	121.000000	42.00
75%	64.500000	80.000000	1.020000	2.000000	0.000000	163.000000	66.00
max	90.000000	180.000000	1.025000	5.000000	5.000000	490.000000	391.00



In [156]: #Finding the Correlation between the plots

```
plt.figure(figsize = (19,19))
sns.heatmap(data.corr(), annot = True, cmap = 'coolwarm') # Looking for strong correlation
```

Out[156]: <AxesSubplot:>



```
In [157]: # checking for null values
```

```
df.isna().sum().sort_values(ascending = False)
```

```
Out[157]: red_blood_cells      152  
red_blood_cell_count     130  
white_blood_cell_count    105  
potassium                  88  
sodium                     87  
packed_cell_volume        70  
pus_cell                   65  
haemoglobin                 52  
sugar                      49  
specific_gravity            47  
albumin                     46  
blood_glucose_random       44  
blood_urea                  19  
serum_creatinine             17  
blood_pressure                12  
age                          9  
bacteria                    4  
pus_cell_clumps              4  
hypertension                  2  
diabetes_mellitus             2  
coronary_artery_disease      2  
appetite                     1  
peda_edema                  1  
aanemia                      1  
class                        0  
dtype: int64
```

```
In [158]: df['packed_cell_volume'] = pd.to_numeric(df['packed_cell_volume'], errors='coerce')  
df['white_blood_cell_count'] = pd.to_numeric(df['white_blood_cell_count'], errors='coerce')  
df['red_blood_cell_count'] = pd.to_numeric(df['red_blood_cell_count'], errors='coerce')
```

```
In [159]: # Extracting categorical and numerical columns
```

```
cat_cols = [col for col in df.columns if df[col].dtype == 'object']  
num_cols = [col for col in df.columns if df[col].dtype != 'object']
```

```
In [160]: # Looking at unique values in categorical columns
for col in cat_cols:
    print(f"{col} has {df[col].unique()} values\n")

red_blood_cells has [nan 'normal' 'abnormal'] values
pus_cell has ['normal' 'abnormal' nan] values
pus_cell_clumps has ['notpresent' 'present' nan] values
bacteria has ['notpresent' 'present' nan] values
hypertension has ['yes' 'no' nan] values
diabetes_mellitus has ['yes' 'no' ' yes' '\tno' '\tyes' nan] values
coronary_artery_disease has ['no' 'yes' '\tno' nan] values
appetite has ['good' 'poor' nan] values
peda_edema has ['no' 'yes' nan] values
aanemia has ['no' 'yes' nan] values
class has ['ckd' 'ckd\t' 'notckd'] values
```

```
In [161]: df['diabetes_mellitus'].replace(to_replace = {'\tno':'no', '\tyes':'yes', ' yes':'yes'}, inplace=True)
df['coronary_artery_disease'] = df['coronary_artery_disease'].replace(to_replace = '\t', value=' ')
df['class'] = df['class'].replace(to_replace = {'ckd\t': 'ckd', 'notckd': 'not ckd'})
```

```
In [162]: df['class'] = df['class'].map({'ckd': 0, 'not ckd': 1})
df['class'] = pd.to_numeric(df['class'], errors='coerce')
```

```
In [163]: cols = ['diabetes_mellitus', 'coronary_artery_disease', 'class']
for col in cols:
    print(f"{col} has {df[col].unique()} values\n")
```

```
diabetes_mellitus has ['yes' 'no' nan] values
coronary_artery_disease has ['no' 'yes' nan] values
class has [0 1] values
```

```
In [164]: df.isna().sum().sort_values(ascending = False)
```

```
Out[164]: red_blood_cells      152
red_blood_cell_count       131
white_blood_cell_count     106
potassium                  88
sodium                     87
packed_cell_volume         71
pus_cell                   65
haemoglobin                52
sugar                      49
specific_gravity           47
albumin                     46
blood_glucose_random       44
blood_urea                  19
serum_creatinine            17
blood_pressure               12
age                          9
bacteria                    4
pus_cell_clumps              4
hypertension                 2
diabetes_mellitus             2
coronary_artery_disease      2
appetite                     1
peda_edema                  1
aanemia                      1
class                        0
dtype: int64
```

```
In [165]: df[num_cols].isnull().sum()
```

```
Out[165]: age                  9
blood_pressure                12
specific_gravity              47
albumin                      46
sugar                         49
blood_glucose_random          44
blood_urea                     19
serum_creatinine               17
sodium                        87
potassium                     88
haemoglobin                   52
packed_cell_volume             71
white_blood_cell_count        106
red_blood_cell_count          131
dtype: int64
```

```
In [166]: df[cat_cols].isnull().sum()
```

```
Out[166]: red_blood_cells      152
          pus_cell              65
          pus_cell_clumps        4
          bacteria               4
          hypertension            2
          diabetes_mellitus       2
          coronary_artery_disease 2
          appetite                1
          peda_edema              1
          aanemia                 1
          class                   0
          dtype: int64
```

```
In [167]: # filling null values, we will use two methods, random sampling for higher null values
# mean/mode sampling for lower null values
```

```
def random_value_imputation(feature):
    random_sample = df[feature].dropna().sample(df[feature].isna().sum())
    random_sample.index = df[df[feature].isnull()].index
    df.loc[df[feature].isnull(), feature] = random_sample

def impute_mode(feature):
    mode = df[feature].mode()[0]
    df[feature] = df[feature].fillna(mode)
```

```
In [168]: # filling num_cols null values using random sampling method
```

```
for col in num_cols:
    random_value_imputation(col)
```

```
In [169]: df[num_cols].isnull().sum()
```

```
Out[169]: age                  0
          blood_pressure        0
          specific_gravity       0
          albumin                0
          sugar                  0
          blood_glucose_random   0
          blood_urea               0
          serum_creatinine        0
          sodium                  0
          potassium               0
          haemoglobin             0
          packed_cell_volume      0
          white_blood_cell_count  0
          red_blood_cell_count    0
          dtype: int64
```

```
In [170]: # filling "red_blood_cells" and "pus_cell" using random sampling method and rest of cat_cols using impute_mode method
```

```
random_value_imputation('red_blood_cells')
random_value_imputation('pus_cell')

for col in cat_cols:
    impute_mode(col)
```

```
In [171]: df[cat_cols].isnull().sum()
```

```
Out[171]: red_blood_cells      0
pus_cell          0
pus_cell_clumps  0
bacteria         0
hypertension      0
diabetes_mellitus 0
coronary_artery_disease 0
appetite          0
peda_edema        0
aanemia           0
class             0
dtype: int64
```

```
In [172]: for col in cat_cols:
    print(f"{col} has {df[col].nunique()} categories\n")
```

```
red_blood_cells has 2 categories
pus_cell has 2 categories
pus_cell_clumps has 2 categories
bacteria has 2 categories
hypertension has 2 categories
diabetes_mellitus has 2 categories
coronary_artery_disease has 2 categories
appetite has 2 categories
peda_edema has 2 categories
aanemia has 2 categories
class has 2 categories
```

```
In [173]: from sklearn.preprocessing import LabelEncoder  
  
le = LabelEncoder()  
  
for col in cat_cols:  
    df[col] = le.fit_transform(df[col])
```

```
In [174]: df.head()
```

Out[174]:

	age	blood_pressure	specific_gravity	albumin	sugar	red_blood_cells	pus_cell	pus_cell_clumps	bact
0	48.0	80.0	1.020	1.0	0.0		1	1	0
1	7.0	50.0	1.020	4.0	0.0		1	1	0
2	62.0	80.0	1.010	2.0	3.0		1	1	0
3	48.0	70.0	1.005	4.0	0.0		1	0	1
4	51.0	80.0	1.010	2.0	0.0		1	1	0



```
In [175]: ind_col = [col for col in df.columns if col != 'class']  
dep_col = 'class'  
X = df[ind_col]  
y = df[dep_col]
```

```
In [176]: # splitting data into training and test set
```

```
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30, random_st
```

K-NEAREST NEIGHBOUR

```
In [177]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

knn = KNeighborsClassifier()
knn.fit(X_train, y_train)

# accuracy score, confusion matrix and classification report of knn

knn_acc = accuracy_score(y_test, knn.predict(X_test))

print(f"Training Accuracy of KNN is {accuracy_score(y_train, knn.predict(X_train))}")
print(f"Test Accuracy of KNN is {knn_acc} \n")

print(f"Confusion Matrix :- \n{confusion_matrix(y_test, knn.predict(X_test))}\n")
print(f"Classification Report :- \n {classification_report(y_test, knn.predict(X_test))}
```

Training Accuracy of KNN is 0.7928571428571428

Test Accuracy of KNN is 0.6

Confusion Matrix :-

```
[[44 28]
 [20 28]]
```

Classification Report :-

	precision	recall	f1-score	support
0	0.69	0.61	0.65	72
1	0.50	0.58	0.54	48
accuracy			0.60	120
macro avg	0.59	0.60	0.59	120
weighted avg	0.61	0.60	0.60	120

DECISION TREE CLASSIFIER

```
In [178]: from sklearn.tree import DecisionTreeClassifier

dtc = DecisionTreeClassifier()
dtc.fit(X_train, y_train)

# accuracy score, confusion matrix and classification report of decision tree

dtc_acc = accuracy_score(y_test, dtc.predict(X_test))

print(f"Training Accuracy of Decision Tree Classifier is {accuracy_score(y_train, dtc)}
print(f"Test Accuracy of Decision Tree Classifier is {dtc_acc} \n")

print(f"Confusion Matrix :- \n{confusion_matrix(y_test, dtc.predict(X_test))}\n")
print(f"Classification Report :- \n {classification_report(y_test, dtc.predict(X_test))}
```

Training Accuracy of Decision Tree Classifier is 1.0
Test Accuracy of Decision Tree Classifier is 0.95

Confusion Matrix :-

```
[[71  1]
 [ 5 43]]
```

Classification Report :-

	precision	recall	f1-score	support
0	0.93	0.99	0.96	72
1	0.98	0.90	0.93	48
accuracy			0.95	120
macro avg	0.96	0.94	0.95	120
weighted avg	0.95	0.95	0.95	120

RANDOM FOREST CLASSIFIER

```
In [179]: from sklearn.ensemble import RandomForestClassifier

rd_clf = RandomForestClassifier(criterion = 'entropy', max_depth = 11, max_features =
rd_clf.fit(X_train, y_train)

# accuracy score, confusion matrix and classification report of random forest

rd_clf_acc = accuracy_score(y_test, rd_clf.predict(X_test))

print(f"Training Accuracy of Random Forest Classifier is {accuracy_score(y_train, rd_c
print(f"Test Accuracy of Random Forest Classifier is {rd_clf_acc} \n")

print(f"Confusion Matrix :- \n{confusion_matrix(y_test, rd_clf.predict(X_test))}\n")
print(f"Classification Report :- \n {classification_report(y_test, rd_clf.predict(X_te
```

Training Accuracy of Random Forest Classifier is 1.0
Test Accuracy of Random Forest Classifier is 0.9583333333333334

Confusion Matrix :-

```
[[72  0]
 [ 5 43]]
```

Classification Report :-

	precision	recall	f1-score	support
0	0.94	1.00	0.97	72
1	1.00	0.90	0.95	48
accuracy			0.96	120
macro avg	0.97	0.95	0.96	120
weighted avg	0.96	0.96	0.96	120

SUPPORT VECTOR REGRESSION

```
In [180]: from sklearn.svm import SVC
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
```

```
In [181]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.15)
```

```
In [183]: SVC = SVC()
print(SVC)

SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

```
-----  
TypeError                                                 Traceback (most recent call last)  
~\AppData\Local\Temp\ipykernel_58256\2678115755.py in <module>  
----> 1 SVC = SVC()  
      2 print(SVC)  
      3  
      4 SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,  
      5       decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',  
  
TypeError: 'SVC' object is not callable
```

```
In [184]: SVC.fit(X_train, y_train)
score = SVR.score(X_train, y_train)
print("Score: ", score)
```

```
-----  
NameError                                                 Traceback (most recent call last)  
~\AppData\Local\Temp\ipykernel_58256\220490471.py in <module>  
      1 SVC.fit(X_train, y_train)  
----> 2 score = SVR.score(X_train, y_train)  
      3 print("Score: ", score)  
  
NameError: name 'SVR' is not defined
```

```
In [185]:  
  
y_pred = SVC.predict(X_test)  
  
cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[46  0]
 [14  0]]
```

LINEAR REGRESSION

```
In [186]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split
import pandas as pd

# Separate the input (X) and output (y) variables
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create a Linear Regression object and fit it to the training data
regressor = LinearRegression()
regressor.fit(X_train, y_train)

# Predict the output values for the testing set
y_pred = regressor.predict(X_test)

# Calculate the R-squared score
accuracy = r2_score(y_test, y_pred)

# Print the accuracy score
print("Accuracy (R-squared):", accuracy)
```

Accuracy (R-squared): 0.7048748428489737

LOGISTIC REGRESSION

```
In [187]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
import pandas as pd

# Load data into a Pandas DataFrame
data = pd.read_csv('kidney_disease.csv')

# Separate the input (X) and output (y) variables
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)

# Create a Logistic Regression object and fit it to the training data
Classifier = LogisticRegression()
Classifier.fit(X_train, y_train)

# Predict the output values for the testing set
y_pred = Classifier.predict(X_test)

# Calculate the accuracy score
accuracy = accuracy_score(y_test, y_pred)

# Print the accuracy score
print("Accuracy:", accuracy)
```

Accuracy: 0.875

RANDOM FOREST

```
In [188]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [189]: X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
```

```
In [190]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=0)
```

```
In [191]: # Instantiate and fit the RandomForestClassifier
forest = RandomForestClassifier()
forest.fit(X_train, y_train)
```

```
Out[191]: RandomForestClassifier()
```

```
In [192]: # Make predictions for the test set
y_pred_test = forest.predict(X_test)
```

```
In [193]: # View accuracy score
accuracy_score(y_test, y_pred_test)
```

Out[193]: 0.95

```
In [194]: # View confusion matrix for test data and predictions
confusion_matrix(y_test, y_pred_test)
```

```
Out[194]: array([[52,  0],
   [ 4, 24]], dtype=int64)
```

```
In [195]: # View the classification report for test data and predictions
print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
0	0.93	1.00	0.96	52
1	1.00	0.86	0.92	28
accuracy			0.95	80
macro avg	0.96	0.93	0.94	80
weighted avg	0.95	0.95	0.95	80

POLYNOMIAL REGRESSION

```
In [196]: from sklearn.preprocessing import PolynomialFeatures
import numpy as np
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=0)
# Transform the input data to include polynomial features up to degree 2
poly = PolynomialFeatures(degree=2)
x_poly = poly.fit_transform(X)
lin_reg_2 = LinearRegression()
lin_reg_2.fit(x_poly, y)
```

Out[196]: LinearRegression()

```
In [197]: lr = LinearRegression()
lr.fit(X_train, y_train)
y_pred = lr.predict(X_test)
print(r2_score(y_test, y_pred))
```

0.7048748428489737

```
In [198]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [199]: X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
```

```
In [200]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=0)
```

```
In [201]: forest = RandomForestClassifier()
forest.fit(X_train, y_train)
```

```
Out[201]: RandomForestClassifier()
```

```
In [202]: y_pred_test = forest.predict(X_test)
```

```
In [203]: accuracy_score(y_test, y_pred_test)
```

```
Out[203]: 0.9625
```

```
In [204]: confusion_matrix(y_test, y_pred_test)
```

```
Out[204]: array([[52,  0],
   [ 3, 25]], dtype=int64)
```

```
In [205]: print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
0	0.95	1.00	0.97	52
1	1.00	0.89	0.94	28
accuracy			0.96	80
macro avg	0.97	0.95	0.96	80
weighted avg	0.96	0.96	0.96	80

Naive bayes

```
In [206]: X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size = 0.2,random_state=0
# training the model on training set
from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb.fit(X_train, y_train)

# making predictions on the testing set
y_pred = gnb.predict(X_test)

# comparing actual response values (y_test) with predicted response values (y_pred)
from sklearn import metrics
print("Gaussian Naive Bayes model accuracy(in %):", metrics.accuracy_score(y_test, y_p
```

Gaussian Naive Bayes model accuracy(in %): 97.5

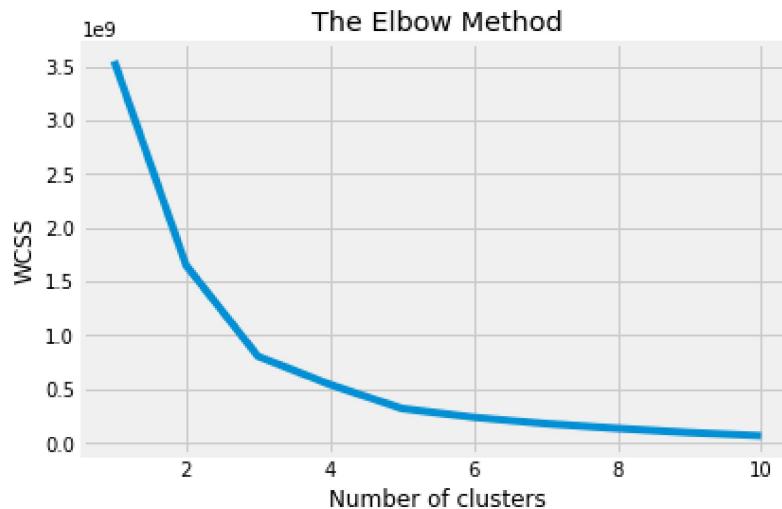
K mean clustering

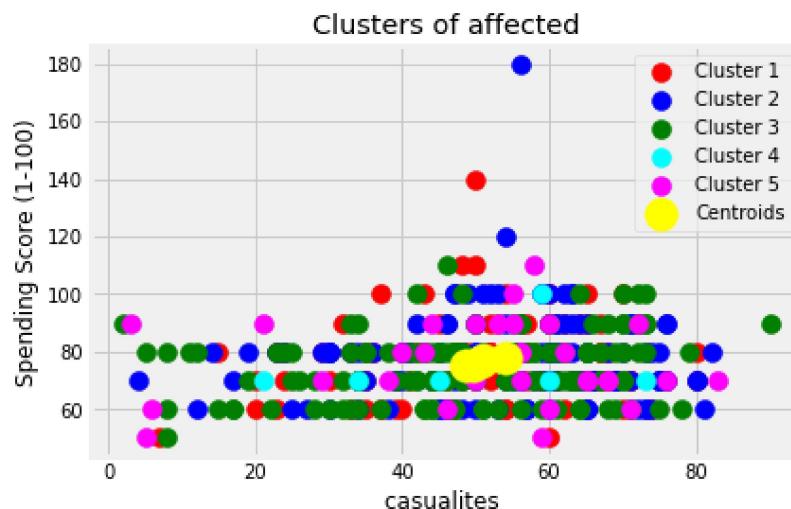
In [83]:

```
from sklearn.cluster import KMeans
wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters = i, init = 'k-means++', random_state = 42)
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)
plt.plot(range(1, 11), wcss)
plt.title('The Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.show()

# Training the K-Means model on the dataset
kmeans = KMeans(n_clusters = 5, init = 'k-means++', random_state = 42)
y_kmeans = kmeans.fit_predict(X)

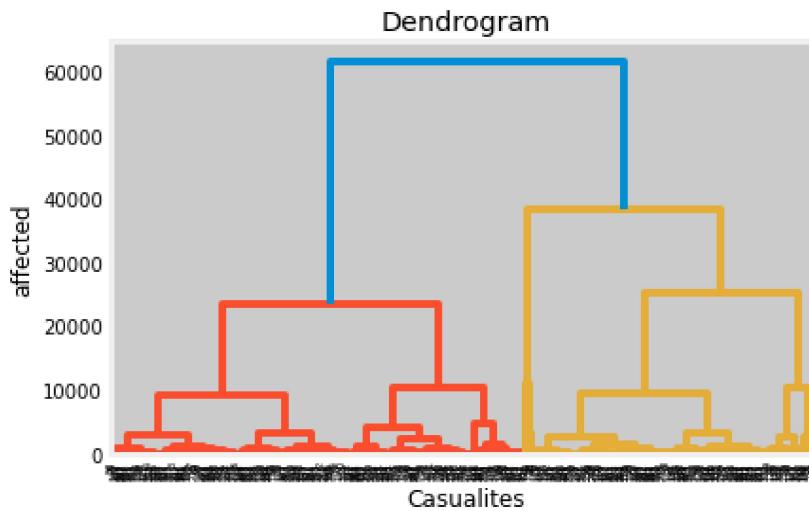
# Visualising the clusters
plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s = 100, c = 'red', label = 'Clu
plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s = 100, c = 'blue', label = 'Cl
plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], s = 100, c = 'green', label = 'C
plt.scatter(X[y_kmeans == 3, 0], X[y_kmeans == 3, 1], s = 100, c = 'cyan', label = 'Cl
plt.scatter(X[y_kmeans == 4, 0], X[y_kmeans == 4, 1], s = 100, c = 'magenta', label =
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1], s = 300, c =
plt.title('Clusters of affected')
plt.xlabel('casualites ')
plt.ylabel('Spending Score (1-100)')
plt.legend()
plt.show()
```





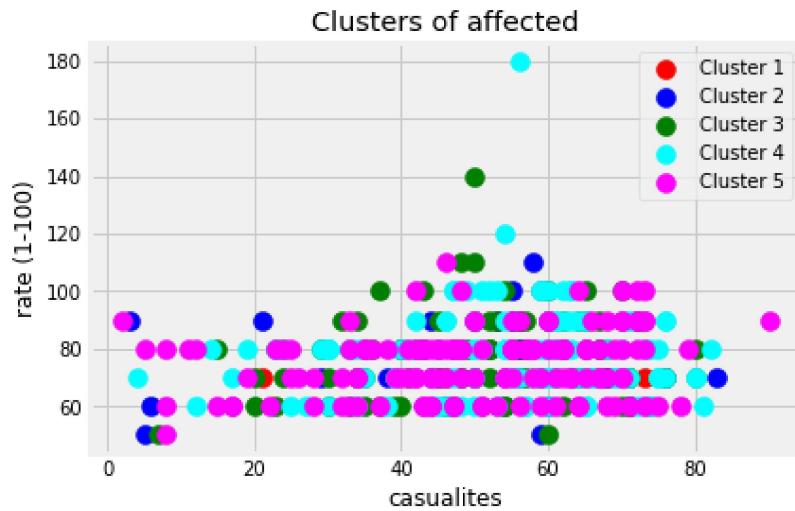
HIERARCHICAL CLUSTERING

```
In [84]: import scipy.cluster.hierarchy as sch
dendrogram = sch.dendrogram(sch.linkage(X, method = 'ward'))
plt.title('Dendrogram')
plt.xlabel('Casualties')
plt.ylabel('affected')
plt.show()
```



```
In [85]: from sklearn.cluster import AgglomerativeClustering
hc = AgglomerativeClustering(n_clusters = 5, affinity = 'euclidean', linkage = 'ward')
y_hc = hc.fit_predict(X)
```

```
In [86]: plt.scatter(X[y_hc == 0, 0], X[y_hc == 0, 1], s = 100, c = 'red', label = 'Cluster 1')
plt.scatter(X[y_hc == 1, 0], X[y_hc == 1, 1], s = 100, c = 'blue', label = 'Cluster 2')
plt.scatter(X[y_hc == 2, 0], X[y_hc == 2, 1], s = 100, c = 'green', label = 'Cluster 3')
plt.scatter(X[y_hc == 3, 0], X[y_hc == 3, 1], s = 100, c = 'cyan', label = 'Cluster 4')
plt.scatter(X[y_hc == 4, 0], X[y_hc == 4, 1], s = 100, c = 'magenta', label = 'Cluster 5')
plt.title('Clusters of affected')
plt.xlabel('casualites')
plt.ylabel('rate (1-100)')
plt.legend()
plt.show()
```



#REINFORCEMENT LEARNING

kernal svm

In [101]:

```
# Import the Libraries
import numpy as np
import matplotlib.pyplot as plt
from sklearn import svm, datasets

X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=0)
from sklearn.svm import SVC
svclassifier = SVC(kernel='rbf')
svclassifier.fit(X_train, y_train)
y_pred = svclassifier.predict(X_test)
from sklearn.metrics import classification_report, confusion_matrix
print("KERNEL SVM WITH GAUSSIAN KERNEL")
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
from sklearn.svm import SVC
svclassifier = SVC(kernel='sigmoid')
svclassifier.fit(X_train, y_train)
y_pred = svclassifier.predict(X_test)
from sklearn.metrics import classification_report, confusion_matrix
print("KERNEL SVM WITH SIGMOID KERNEL")
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

KERNEL SVM WITH GAUSSIAN KERNEL

```
[[52  0]
 [28  0]]
```

	precision	recall	f1-score	support
0	0.65	1.00	0.79	52
1	0.00	0.00	0.00	28
accuracy			0.65	80
macro avg	0.33	0.50	0.39	80
weighted avg	0.42	0.65	0.51	80

KERNEL SVM WITH SIGMOID KERNEL

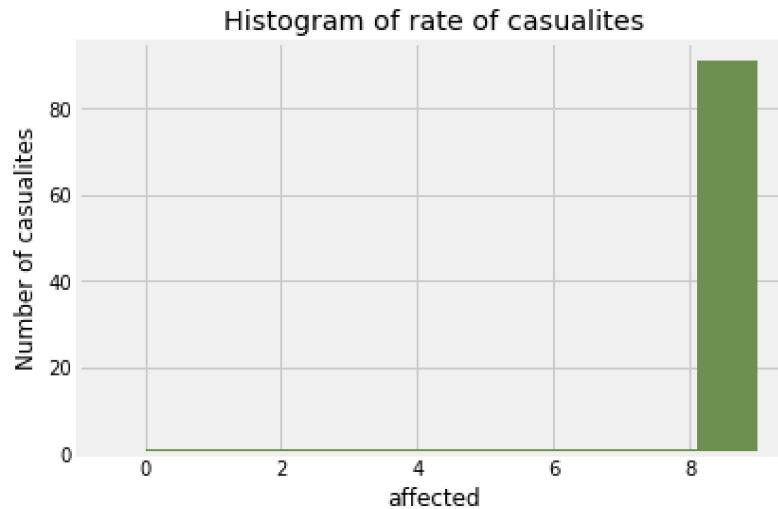
```
[[28 24]
 [25  3]]
```

	precision	recall	f1-score	support
0	0.53	0.54	0.53	52
1	0.11	0.11	0.11	28
accuracy			0.39	80
macro avg	0.32	0.32	0.32	80
weighted avg	0.38	0.39	0.38	80

Upper confidence bound

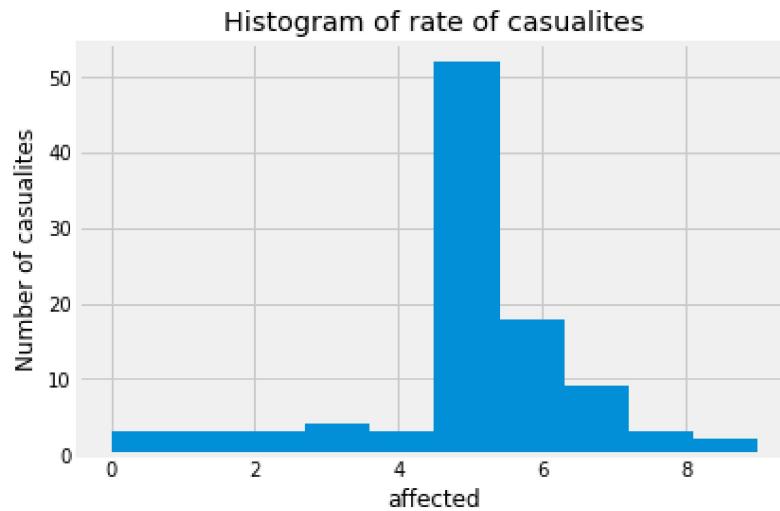
In [207]:

```
# Implementing UCB
import math
N = 100
d = 10
ads_selected = []
numbers_of_selections = [0] * d
sums_of_rewards = [0] * d
total_reward = 0
for n in range(0, N):
    ad = 0
    max_upper_bound = 0
    for i in range(0, d):
        if (numbers_of_selections[i] > 0):
            average_reward = sums_of_rewards[i] / numbers_of_selections[i]
            delta_i = math.sqrt(3/2 * math.log(n + 1) / numbers_of_selections[i])
            upper_bound = average_reward + delta_i
        else:
            upper_bound = 1e200
        if upper_bound > max_upper_bound:
            max_upper_bound = upper_bound
            ad = i
    ads_selected.append(ad)
    numbers_of_selections[ad] = numbers_of_selections[ad] + 1
    reward = df.values[n, ad]
    sums_of_rewards[ad] = sums_of_rewards[ad] + reward
    total_reward = total_reward + reward
    plt.hist(ads_selected)
plt.title('Histogram of rate of casualites')
plt.xlabel('affected')
plt.ylabel('Number of casualites')
plt.show()
```



```
In [129]: #thomson bound
import random
N = 100
d = 10
ads_selected = []
numbers_of_rewards_1 = [0] * d
numbers_of_rewards_0 = [0] * d
total_reward = 0
for n in range(0, N):
    ad = 0
    max_random = 0
    for i in range(0, d):
        random_beta = random.betavariate(numbers_of_rewards_1[i] + 1, numbers_of_rewards_0[i] + 1)
        if random_beta > max_random:
            max_random = random_beta
            ad = i
    ads_selected.append(ad)
    reward = df.values[n, ad]
    if reward == 1:
        numbers_of_rewards_1[ad] = numbers_of_rewards_1[ad] + 1
    else:
        numbers_of_rewards_0[ad] = numbers_of_rewards_0[ad] + 1
    total_reward = total_reward + reward

# Visualising the results - Histogram
plt.hist(ads_selected)
plt.title('Histogram of rate of casualites')
plt.xlabel('affected')
plt.ylabel('Number of casualites')
plt.show()
```



```
In [ ]:
```